

The Prompt Ladder

Aditya Dixit | FlyRank AI Internship — General AI Fluency, Week 2 (Foundations) | Track: Frontend AI Engineering

• Assignment:** CUSTOM-MRBSNHC-B3701520 | Workload: 2h | Phase: Foundations

The cheapest performance upgrade in AI is not a new model — it's changing one thing at a time and comparing outputs. This ladder does exactly that.

Baseline (the embarrassing prompt I would have typed a month ago)

• Prompt — Baseline (0):**

Build a website

That's it. No goal, no audience, no context, no format, no constraints. The kind of prompt that feels fast and teaches nothing.

• Output — Baseline (excerpt, as returned):**

`<html`

My Website

Welcome to My Website

Lorem ipsum dolor sit amet, consectetur adipiscing elit...

Click Here

...

*Generic HTML, lorem ipsum, "Click Here" CTA, no Tailwind, no React, no real business. It could be a dentist, a DJ, or a dog food store. Zero signal about what the site should *do*.*

Version 1 — Layer added: **clearer goal**

• Prompt — V1:**

Build a single-page marketing website that gets visitors to book a coffee tasting. Keep the page focused on one action.

• What changed in prompt:** Added one layer — a clearer goal (book a tasting, not "build a website").

• Output — V1 (excerpt):**

Hero: "Experience Coffee Differently" + CTA "Book Your Tasting"

Sections: Hero → About → Menu → Contact (generic filler copy)

Still HTML boilerplate with placeholder text: "Our coffee is the best in town..."

• Four notes — V1:**

• **What changed in prompt:** Added clearer goal (one action: book tasting).

• **What actually improved in the output:** It stopped generating a random multi-purpose site and finally anchored on a single CTA — the hero now has "Book Your Tasting" instead of "Click Here," so a visitor knows what to do.

- **What still failed:** Copy is still generic ("best in town"), no audience, no real shop details, layout is a stock 4-section template that could be any café.
- **What I'd try next:** Define *who* I'm talking to — busy professionals vs students changes everything about tone and what to emphasize.

Version 2 — Layer added: **defined audience**

- Prompt — V2:

Build a single-page marketing website that gets busy professionals in Kanpur who know coffee but are new to specialty roasters to book a coffee tasting.

- What changed in prompt: Added one layer — defined audience (Kanpur professionals, coffee-aware but specialty-new).

- Output — V2 (excerpt):

Hero: "For Professionals Who Take Coffee Seriously"

Copy: "Elevate your mornings with our premium single-origin blends, curated for the discerning connoisseur..."

Tone shifted to corporate: "Optimize your productivity with artisanal excellence"

- Four notes — V2:

- **What changed in prompt:** Added defined audience.
- **What actually improved in the output:** *This one didn't help — it made it worse.* The output started using jargon my audience wouldn't say ("discerning connoisseur," "optimize productivity") and assumed a stiff corporate reader. It sounded less like a neighbourhood roaster and more like a LinkedIn ad. The lesson: audience without real context just invites stereotypes.
- **What still failed:** Still no real shop, no menu, no location, no palette — so the model filled the gap with buzzwords.
- **What I'd try next:** Give it real context — my actual shop, my actual product — so it stops inventing a fake audience persona.

Version 3 — Layer added: **real context**

- Prompt — V3:

Build a single-page marketing website that gets busy professionals in Kanpur who know coffee but are new to specialty roasters to book a coffee tasting.

Context: This is for Brew & Co — my specialty coffee roaster template (React + Vite + Tailwind + Framer Motion, live at brew-and-co-opal.vercel.app). Light-roast focus, direct-trade story, warm minimal palette (#f172a / #f8fac / amber accents), existing sections: hero, story, roast lineup, tasting booking. Use real Kanpur details, not lorem ipsum.

- What changed in prompt: Added one layer — real context (actual project, stack, palette, story, live URL).

- Output — V3 (excerpt):

Hero: "Light Roast, Direct Trade — Tasted in Kanpur"

Story: "We source from Chikmagalur estate, roast light to keep floral notes..."

Roast Lineup: "Ethiopia Guji (jasmine), Chikmagalur Honey (stone fruit)" with tasting notes

CTA: "Book a 20-min tasting — pick your slot, taste two pours side by side"

- Four notes — V3:

- **What changed in prompt:** Added real context (Brew & Co's actual roasts, story, stack, palette).

- **What actually improved in the output:** It stopped inventing fake coffee and finally sounded like *my* shop — real origins, real tasting notes, real Kanpur anchor. A stranger can now tell this isn't a template; it's a specific roaster with a point of view. Biggest jump so far.
- **What still failed:** Output came back as one giant markdown dump with mixed HTML/Tailwind snippets and no file structure — I couldn't copy-paste it into my Vite project without hunting through prose.
- **What I'd try next:** Specify the output format so the code is actually usable.

Version 4 — Layer added: **specified output format**

- Prompt — V4:

Build a single-page marketing website that gets busy professionals in Kanpur who know coffee but are new to specialty roasters to book a coffee tasting.

Context: Brew & Co — React + Vite + Tailwind + Framer Motion, light-roast direct-trade, palette #0f172a/#f8fac/amber, sections hero/story/lineup/booking, live at brew-and-co-opal.vercel.app.

Output format: 1) File map (App.jsx, components/Hero.jsx etc.), 2) Full code for Hero + Booking section only (React + Tailwind, no lorem ipsum, copy in my voice: direct, warm, plain, specific), 3) Brief notes on where placeholder images go. No extra prose.

- What changed in prompt: Added one layer — specified output format (file map + scoped code blocks + voice).

- Output — V4 (excerpt):

File Map:

- src/App.jsx

- src/components/Hero.jsx

- src/components/Booking.jsx

```jsx

// Hero.jsx

Taste light roast, properly — in Kanpur

Two pours, one story. Direct trade from Chikmagalur, roasted light.

Book a 20-min tasting

...

(followed by 80 more lines of Booking form code)

- Four notes — V4:

- **What changed in prompt:** Added specified output format.
 - **What actually improved in the output:** It became copy-paste usable — I got a file map and real React/Tailwind code for Hero + Booking instead of a wall of prose. I could drop it into `src/` immediately.
 - **What still failed:** *Honest "made it a bit worse" moment:* it became *too* verbose — it generated full props, Framer variants, and 3 example bookings even though I asked for 2 sections, and the copy repeated "Kanpur" 6 times because the format didn't constrain length. Useful but not tight.
 - **What I'd try next:** Add constraints so it stays runnable and tight.
-

Version 5 — Layer added: ****constraints****

• Prompt — V5:**

Build a single-page marketing website that gets busy professionals in Kanpur who know coffee but are new to specialty roasters to book a coffee tasting.

Context: Brew & Co — React + Vite + Tailwind + Framer Motion, light-roast direct-trade, palette #0f172a/#f8fafc/amber, sections hero/story/lineup/booking.

Output format: File map + full React + Tailwind code for Hero + Booking only, copy in voice: direct, warm, plain, specific, no buzzwords.

Constraints: mobile-first, Tailwind only (no extra deps beyond React/Framer), accessible (labels, aria, keyboard), max 18 words per headline, no lorem ipsum, must be copy-paste runnable. Before final, self-check: does npm run build pass? Is the booking button reachable by keyboard? Fix if not.

• What changed in prompt:** Added one layer — constraints (mobile-first, stack-locked, a11y, length, runnable check).

• Output — V5 (excerpt):**

```
``jsx

// Hero.jsx — 22 lines, mobile-first

Light roast, direct trade — taste it here

20 minutes, two pours. Same bean, two roasts — you pick what you like.

Book a tasting — 20 min


// Booking.jsx — labels, aria-describedby, keyboard focus ring, no extra deps
...

Self-check note: "Build passes, button has focus ring, headline 7 words, no lorem ipsum."
```

• Four notes — V5:**

- ****What changed in prompt:**** Added constraints (mobile-first, stack, a11y, length, self-check).
- ****What actually improved in the output:**** It got tight and shippable — headline dropped from 12 words to 7, mobile padding fixed, button got a focus ring, and it stopped adding random dependencies. The self-check actually caught a missing ``htmlFor`` and fixed it before I had to.
- ****What still failed:**** Self-check is honest but still self-reported — the model can claim "build passes" without running it. I still need a real verification step when I paste it.
- ****What I'd try next for a stranger:**** Bundle the best layers, add a verification requirement ("paste, run build, list 3 checks you passed"), and clean placeholders so someone else can run it without me.

Side-by-side comparison (don't trust memory)

Version Layer added What the output *felt* like		
----- ----- -----		
Baseline — Generic dentist-site HTML		
V1 clearer goal Finally a CTA, still lorem ipsum		
V2 audience (alone) Worse — corporate jargon, fake persona		

| V3 | real context | Huge leap — real beans, real story |

| V4 | format | Usable code, but verbose/repetitive |

| V5 | constraints | Tight, accessible, runnable |

Final reusable prompt (works for a stranger on Frontend AI Engineering)

Paste this as-is — swap the bracketed bits. No bundling confusion, no "you need to know my project" gap.

*You are building a single-page marketing site in **React + Vite + Tailwind + Framer Motion** (no other deps) that gets **[AUDIENCE — e.g., busy professionals in Kanpur who know coffee but are new to specialty]** to take **[ONE GOAL — e.g., book a 20-min coffee tasting]**.*

***Real context (replace with yours):** Project = **[PROJECT NAME + one-line story — e.g., Brew & Co, light-roast direct-trade roaster, Chikmagalur beans, warm minimal palette #0f172a/#f8f4e6/amber]**. Existing stack/live URL: **[e.g., brew-and-co-opal.vercel.app]**. Sections already assumed: hero, story, lineup, booking.*

Output format (exactly this):

1) File map (e.g., App.jsx, components/Hero.jsx, components/Booking.jsx)

*2) Full code for **Hero + Booking only** (React + Tailwind, mobile-first, copy-paste runnable)*

3) One-line note per placeholder image

***Constraints:** Mobile-first, Tailwind only, accessible (labels + htmlFor/id, aria-invalid/describedby where needed, keyboard focus visible), max 18 words per headline, no lorem ipsum, copy in voice: **direct, warm, plain, specific, no buzzwords**.*

***Verification before you finish:** After the code, add a 3-line self-check: (a) would `npm run build` pass? (b) is the primary CTA reachable by keyboard with visible focus? (c) is every headline ≤18 words? Fix any fail before final.*

• Why this works for a stranger: It names the audience and goal up front, forces real context (so no invented beans), locks format (so code is paste-ready), constrains stack/a11y/length (so output stays tight), and requires verification (so the model can't just claim "done").

• eliverable checklist: 6 runs (baseline + 5) ✓ | Each version = exactly one named layer ✓ | Notes describe output changes, not just prompt changes ✓ | At least one honest "didn't help / made it worse" (V2 audience alone, V4 verbosity) ✓ | Final prompt works without me ✓